

<https://www.kaggle.com/datasets/retailrocket/ecommerce-dataset?resource=download>

Technical Documentation

Project: Product Embedding-based Recommendation System

Overview

This document describes the end-to-end architecture, data preparation, embedding neural network design, training, and serving strategy of a product recommendation system based on embedding techniques. The embeddings are learned from sequences of product views by users, leveraging the assumption that products viewed in similar contexts have similar characteristics.

1. Dataset

Source:

- **Retail Rocket eCommerce Dataset (Kaggle)**
 - events.csv: timestamped product view data
 - Columns: timestamp, visitorid, event, itemid

Raw Data Sample:

timestamp	visitorid	event	itemid
2015-06-02 05:02:12	257597	view	355908
2015-06-02 05:50:14	992329	view	248676
...

2. Data Preparation Pipeline

a. Sessionization:

- Group sequential product views into user sessions based on inactivity (e.g., 30 minutes).
- **Output:** Sessions as ordered product-view sequences:

SessionID 123 → [355908, 248676, 318965]

b. Generating Training Pairs (Positive samples):

- For every session, generate pairs of viewed products with a window size = 1:

[(248676, 355908), (318965, 248676), ...]

- Here, (target, context) implies context is viewed right before or after target.

c. Negative Sampling:

- For each positive pair (target, context), randomly sample products that did not appear with the target in that session as negative samples.
- **Output:** Final labeled training set:

[(248676, 355908, 1), (248676, 111111, 0), (318965, 248676, 1), (318965, 222222, 0)]

d. Mapping item_id to Integer Indices:

- Embedding layers require dense integer IDs:

item2idx = {248676:0, 318965:1, 355908:2, ...}

idx2item = {0:248676, 1:318965, 2:355908, ...}

3. Neural Network Architecture

a. Model Inputs:

- Two input IDs (target, context), both integers representing products.

b. Embedding Layers:

- Separate embedding layers for target and context:

target_embedding $\in \mathbb{R}^{(N_items \times d)}$

context_embedding $\in \mathbb{R}^{(N_items \times d)}$

- N_items = Total unique products.
- d = Embedding dimension (e.g., 50).

c. Forward Computation:

For single training sample (t, c) pair:

$v_t = \text{target_embedding}[t] \in \mathbb{R}^{(1 \times d)}$

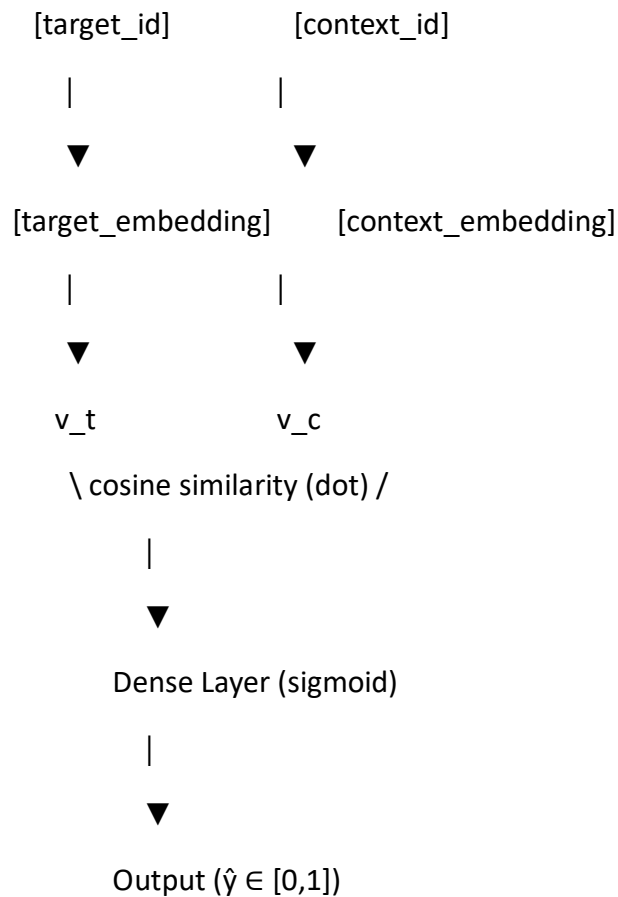
$v_c = \text{context_embedding}[c] \in \mathbb{R}^{(1 \times d)}$

$\text{cos_sim} = \text{cosine_similarity}(v_t, v_c)$ # normalized dot-product

$\text{dense_output} = \text{sigmoid}(w \cdot \text{cos_sim} + b)$

- Output probability (dense_output) predicts whether the given pair (t, c) is legitimate (viewed together).

Architecture Diagram:



4. Training Procedure (Detailed)

a. Loss Function:

- Binary Cross Entropy (BCE):

$$L(y, \hat{y}) = -(y \log(\hat{y}) + (1-y) \log(1-\hat{y})) \quad L(y, \hat{y}) = -(y \log(\hat{y}) + (1-y) \log(1-\hat{y}))$$

b. Optimization:

- Optimizer: Adam
- Batch training with mini-batches (e.g., 512 pairs/batch).

c. Parameter Updates:

- Only the rows of embeddings corresponding to the target and context product indices in the current batch are updated.

$$W_{tgt}[t] \leftarrow W_{tgt}[t] - \eta \nabla(W_{tgt}[t])$$

$$W_{ctx}[c] \leftarrow W_{ctx}[c] - \eta \nabla(W_{ctx}[c])$$

d. Final Trained Embedding Matrix:

- After training, embeddings are extracted and normalized to create product vectors for serving recommendations.

5. Serving / Deployment Architecture

a. Preparation (Offline):

- Embedding matrix E is row-wise L2-normalized for fast cosine similarity computations:

$$E_{normalized}[i] = E[i] / ||E[i]||$$

- Embeddings stored and indexed using Approximate Nearest Neighbor (ANN) search (e.g., FAISS, Annoy).

b. Online Recommendation Workflow:

User views product X:

↳ Fetch embedding vector V_x from $E_{normalized}$.

↳ Perform ANN lookup to find top-N nearest vectors.

↳ Map nearest vectors' indices back to real product IDs.

↳ Business filtering (stock, promotions, etc.).

↳ Return recommendations to frontend.

6. Evaluation & Validation:

- **Quantitative:**
 - Classification accuracy, precision-recall for the binary prediction task.
 - Offline metrics: Hit Rate, NDCG, MRR.
 - **Qualitative:**
 - Inspect embedding quality by visualizing using T-SNE or UMAP.
-

7. Low-Level Technical Specifications:

Frameworks & Libraries:

- Python, Pandas, NumPy
- TensorFlow / Keras / PyTorch for neural networks
- FAISS/Annoy for ANN search

Model Serving:

- Embeddings loaded into memory or ANN indices.
 - RESTful API (FastAPI, Flask, or similar) for serving requests.
-

8. Limitations & Extensions:

Limitations:

- Cold-start problem for new products without embedding vectors.
- Does not inherently capture user personalization.

Possible Extensions:

- Integrate user embeddings to personalize recommendations.
 - Combine metadata embeddings (images, text) with product embeddings.
 - Periodically retrain embeddings to adapt to product catalog changes.
-

9. Summary of the Flow:

Raw events → Sessionization → Positive/Negative Sampling → Integer Mapping →

Neural Network Training → Extract Embeddings → Normalize → ANN Indexing → Serve Recommendations

Conclusion

This document thoroughly captures the data-flow, computations, technical aspects, and low-level architecture necessary to implement a scalable, efficient embedding-based recommender system. This architecture ensures recommendations leverage implicit user interaction signals, significantly enhancing user experience and product discovery.
